Practical 1 - *Know your Data & Predictions*

Chris Emmery

2 September 2017

In this practical[1] we will focus (i) on exploratory dataset analysis - an important first step before engaging in any Data Mining activities. We want get to know our data: find interesting observations, anomalies in data, and think about how to best go about solving the task. After, (ii) on making predictions, ascertaining the correctness of these predictions, and trying to party improve them.

Refresher

DATA MINING, MACHINE LEARNING, AND ALL RELATED FIELDS have a special way of naming certain things related to data that are important to know by heart. This refresher will give you a short overview. Say we have a table (our data) describing ticket sales for concerts. See a small example below:

id	gender	age	price	concert	genre
01	f	21	40	AC/DC	hard rock
07	m	45	75	Ed Sheeran	acoustic
11	m	30	64	Depeche Mode	electronic
05	f	14	54	Hans Zimmer	instrumental

Each line (or row) of data is what we call an **instance**[2]. An instance usually refers to something concrete: a person, an animal, a song, a document, or in this case: a ticket. Instances are described by **features**[3], they tell you some (hopefully) characteristic information about these instances. Instances can be represented mathematically as a feature vector[4].

Interpreting Features: Continuous vs Discrete

THIS DOESN'T FULLY WORK, HOWEVER. In general, we want to be able to do calculations with these data. As such, features with what we call strings (letters) cannot be used in their raw form (e.g. `f`, `AC/DC`). For the sake of simplicity, we will get rid of the concert and genre features for this example. The gender variable can be recoded to 0 for f and 1 for m. So, we will get:

$$\vec{x}_1 = \langle 01, 0, 21, 40 \rangle \tag{2}$$

[1] **Important Practical Note**: If you cannot answer a task during the practicals fully, or feel unsure about your answer (even after the explanation), please ask! It is very important that you develop the correct intuitions for each of the points we discuss here. Sometimes they just don't 'click' by themselves; they require a lot of repeated practice and interpretation, and not every explanation works for everyone. We will be very happy to answer all your questions on the Forum or during class!

Table 1: Tickets instances containing customer & meta-data regarding their sales.

[2] Can also be referred to as an entry, observation, or data point.

[3] Also known as variables, or attributes.

[4] For the first instance above, this would be noted as:

$$\vec{x}_1 = \langle 01, f, 21, 40, AC/DC, hard\ rock \rangle \tag{1}$$

Where $\vec{x}$ is the vector, and subscript 1 the index.

Figure 1: A 3d representation where classes are coloured.

Now, the remaining 4 features also require a bit of interpretation. We distinguish between **continuous** features and **discrete** features. Continuous features are real valued and can be within some range, whereas discrete features are finite, and are usually associated with some label or category. In our example, `id`, `age` and `price` are all considered numeric (technically, they are all **ratio**). Gender can either be **binary** feature (0 = f and 1 = m), or other gender identities might be added, making it **nominal**. See Figure 2 for further info.

Feature Spaces

THE BEAUTY OF USING THIS NUMERIC REPRESENTATION is that each instance can be represented as a point in a space (or n-dimensional graph). So say that we plot the age on the x and price on the y-axis for some instances, our feature space as this is called, will look like Figure 3. This is a 2-dimensional space (with 2 axes). A bit more fancier 3-d space can be in Figure 1. However, we have 4 features - so this would have to be plotted in a 4-dimensional space. Unfortunately, we as humans can only see as much as 3-dimensions, so there's no way to show you how our actual feature space would look like (even colors and shapes will only get us so far). Luckily however, mathematical functions can handle this!

Part A - Understanding your Data

YOUR FIRST STEP AS A DATA SCIENTIST is (as frequently repeated in the lecture) to *know* your data. Everyone can learn to fire up a program and click a few buttons, or write a few lines of code. It is your task, however, to understand this data and generate creative insights and be able to communicate these (in the form of a scientific paper or presentation). It is therefore important to know the possibilities and limitations of your data. The first dataset we'll be working with is the following:

Information	task	data
Data	raw	github

← Text in the table is clickable.

Please make sure you understand the dataset and the task **before** beginning!

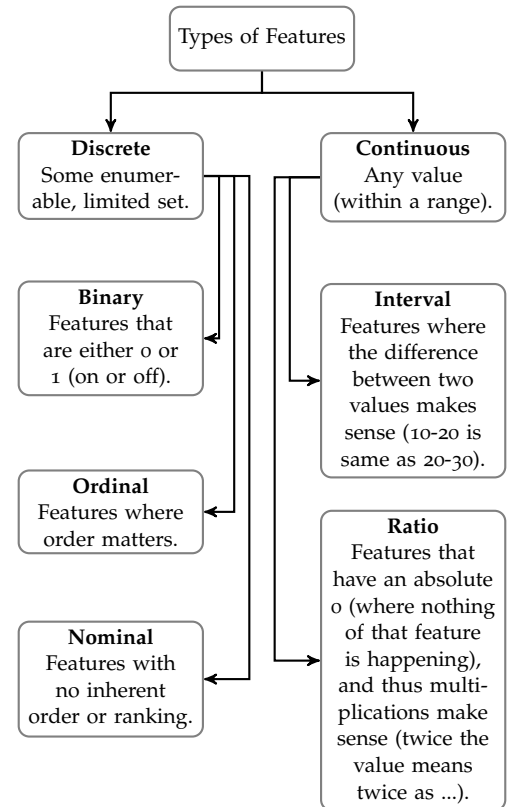Figure 2: Overview of the different types of features.

Figure 3: Vector space representation of Table 1.

Setting up Jupyter

Open up Jupyter and do the following:

- Sync the `data-mining` repository by running `!git pull` in your main Jupyter directory (see previous video).

- Please put your ANR in the first cell. As shown below.

Start with inserting ANR in the first cell (see Figure 4), then switch the cell from `code` to `markdown` (for typing) by switching in the toolbar (Figure 5). Finally, you can insert new cells below using the *Insert* part of the menu (Figure 6). Alternatively, while having your cursor active in a cell, press `Esc` to deselect it, and then `b`. You can either run a code cell, or render a Markdown[5] cell, by pressing `ctrl`/`cmd` + `Enter`. For other shortcuts, see `Help` `Keyboard Shortcuts`.

Now, let's get started, we're first going to create a cell with a few handy tools for the notebook. The first line `%matplotlib inline` makes sure that whenever there's a plot, it is displayed in our browser (otherwise it would just output some text). The others use the 'warning' package and disable some warnings from Python (pandas sometimes throws warnings for functions that will not work in some new version and such). These clutter our notebook, so we don't want them[6].

```
%matplotlib inline

import warnings
warnings.simplefilter(action='ignore')
```

This is where we got last week, but we loaded `housing.csv`. Now, we want to load `titanic.csv`. A small explanation of the directory: `.` means 'from the current directory'. If you check the repository, you can see that `data` is in the main directory (`data-mining`). If you create the notebook one directory further down (say in Week 2), you would specify `../data/Titanic/titanic.csv`, where `..` means 'one directory back'. The Titanic dataset has an actual index column, namely `PassengerId`. This is unique for every instance, and we don't want to use it as a feature anyway (think about why).

We shorthand pandas (version 0.14.1) to `pd` and then load the `.csv` file into our `DataFrame` using `from_csv`. I will be providing links to the documentation of the objects and functions that we're using. Here you can see information regarding the **parameters** that a function takes. They are basically 'options' for a certain function. As you can see, `from_csv` has `index_col`, but also for example `sep`, which—if you read the documentation—can be used to indicate the `.csv` file


Figure 4: ANR: ....

Figure 5: Switching to Markdown.

Figure 6: Insert new cell.

[5] Markdown is actually a pretty handy standard for formatting text. It's easy to practice within Jupyter.

[6] **Important Notes**: `File` `Save and Checkpoint` often! The servers can be unstable. This prevents wasted time and (potentially) headaches. Also, in the code it often says ', e.g. around 'ignore' here. These are supposed to be **single quotes**: the typewriter apostrophe. Don't copy paste them, or you'll get a `SyntaxError: invalid character in identifier`!

delimiter. 'Delimiter' here means the character separating the values in our `.csv` file (by default it's set to `','` which is according to the format of our file).

Finally we store our dataframe in what's called a **variable**; a container with a name that we store our data (or any other object) under. Here, I call it `df` for dataframe. If we put this variable name on the last line in Jupyter, it will output the contents of the variable (in this case, our data). As such[7]:

```python
import pandas as pd

df = pd.DataFrame.from_csv('./data/Titanic/titanic.csv',
                          index_col='PassengerId')
df
```

Now let's get into some of the operations that we can run on our `DataFrame`. If you search the DataFrame documentation for `Methods`, they are all listed. These are all click-able, where more of their functionality is explained. It also short-lists the parameters of the function, if they are in between brackets (`[like,so]`), they are optional. The first one we'll be using is `count()` (that, as you can see in the docs, only has optional parameters). Let's try:

```python
df.count()
```

When you put a function on the last line of a cell without a variable to store it in, this will make Jupyter also show the output of the function (like above).

Interpreting Raw Data

GIVEN THIS NEW INFORMATION try solving the following tasks. If we take a quick peek at the Wiki (link), we get some rough numbers for comparison[8].

Task 1

- How much of the data are we missing?

- Do you see any features that you do not understand the values of?

- Which information could we potentially use to determine the crew of the Titanic in this dataset?

Solutions for all tasks are listed at the end.

[7] Please keep in mind that each dark block of code is supposed to be a new cell! Read back up if you forgot how to create those.

[8] "Titanic had around 885 crew members on board for her maiden voyage.[95] Like other vessels of her time, she did not have a permanent crew, and the vast majority of crew members were casual workers who only came aboard the ship a few hours before she sailed from Southampton.[96]

Titanic's passengers numbered approximately 1,317 people: 324 in First Class, 284 in Second Class, and 709 in Third Class. Of these, 869 (66%) were male and 447 (34%) female. There were 107 children aboard, the largest number of which were in Third Class.[105] The ship was considerably under capacity on her maiden voyage, as she could accommodate 2,453 passengers—833 First Class, 614 Second Class, and 1,006 Third Class."

We can try practising parameter-use a bit by trying the `axis` parameter to get column-wise counts rather than row-wise. As you can see from this line in the docs:

```
DataFrame.count(axis=0,level=None, numeric_only=False)
```

Axis is set to `0` by default.

```
df.count(axis=1)
```

Solve the difference between row-wise and column-wise counts for yourself. Also note that when getting any sort of counts for rows, they are always listed together with their index (`PassengerId`). We can see which passenger paid how much fare by selecting a specific column. This is done with:

```
df['Fare']
```

If we'd like to add the names to the output, we can list them like so:

```
df[['Name', 'Fare']]
```

Let's try aggregating some data. We're interested in how many passengers were in which economic class. Note that when subsetting our `pd.DataFrame` object to only a single column, it actually turns into a `Series`. These have different methods (functions) than Dataframes. Try `value_counts()`:

```
df['Pclass'].value_counts()
```

We can also aggregate even more 'complex' levels. Let's try grouping by Sex using `groupby()`, which we simply add in between our previous line:

```
df.groupby('Sex')['Pclass'].value_counts()
```

Just to quickly walk you through the steps: `groupby` splits the entire dataframe between sexes (so all feature values get divided between male and female). Convince yourself of this by running:

```
df.groupby('Sex').count()
```

Notice that `count` is used here again, because `groupby` returns all columns in the Dataframe, so the object is still a Dataframe (not a Series). This in contrast to the line before, where we add `['Pclass']` to only select the class feature (turning it into a Series, because only

one column). For that one, we'd apply the `value_counts()`, as you can see. So, important take-away: always think if you're working on more than one column (dataframe), or just one (series), and see which methods you can use accordingly. Now we can safely go into exploring more methods.

Visualization

PANDAS WORKS WELL for selecting (sub-setting) and plotting different entries within a dataset. We'll first look at a way to visualize certain feature distributions using **density plots** (see Figure 7, and see if we can group them by `Survived`. For all the features that we inspect, you will see if this influenced whether people survived the Titanic yes or no. Say that we'd want to check this for `Fare`, the first step would be:

```
df.groupby('Survived')['Fare']
```

Both DataFrames and Series have a `plot()` method, with many possible parameters. You're free to check them out, for this we'll be using KDE (Kernel Density Estimation, or density plot). We make our line look like this:

```
df.groupby('Survived')['Fare'].plot(kind='kde', legend=True)
```

The default plots in Pandas don't look too fancy (granted, it's not a plotting library, this is just convenient functionality). We can alter this somewhat by for example:

```
import matplotlib.pyplot as plt
plt.style.use('ggplot')
```

Now try running the plot again! Now, if you start changing around the features, you'll notice that some aren't compatible with density plots (they throw an error). This is because they are discrete, and therefore don't have a density estimation. You can use bar plots for these:

```
df.groupby('Survived')['Sex'].value_counts().plot(kind='Bar')
```

You should be able to figure out why we didn't include `legend` here. If you tried age, you see the warning `array must not contain infs or NaNs`. If you look at the very first `pd` table dump, you'll notice


Figure 7: Density plot.

there are NaN ages (Not a Number). These are missing values. We can fix this with overwriting the age feature:

```
df['Age'] = df['Age'].fillna(0)
```

Here, we fill the NaNs with 0. Now, when interpreting the results, take into account that zeroes are actually NaNs. You can up the complexity of the groupby function by stacking more features as a list, like so:

```
df.groupby(['Sex', 'Pclass'])['Survived'].value_counts().plot(kind='Bar')
```

You can also show the counts above in a table when removing `.plot(kind='Bar')`, if you replace `value_counts()` with `mean()` you get the averages for each group split, etc. Using these plots and checking the table outputs you can answer the following questions:

Task 2

- Find two features where a certain group has a noticeable high mortality rate.

- Explain why this is the case with your knowledge of the dataset (having seen the movie might help, see Figure 8).

Distributions are pretty straight-forward to interpret; the visualization between two features is simple and effective. However, it also limits displaying interaction between several features (e.g. if young males tended have lower survival rates than females). For this, we can use **scatter plots** (see Figure 9). In pandas, you need to set an *x* and *y* label, like so:

```
df.plot(x='somefeature', y='someotherfeature', kind='scatter')
```

To stack multiple subsets in one plot, and give them different colors, we have to select them from the DataFrame based on some condition. We can do this for example for `Survived` with the following piece of code:

```
died_subset = df[df['Survived'] == 0]
live_subset = df[df['Survived'] != 0]
```

To break this down, the inner part (`df['Survived'] == 0`) returns the index numbers (that's where they come in handy!) where the value of `Survived` was equal to (==) zero (so the people who did not survive). The outer part `df[....]` then receives those index numbers,


Figure 8: Hint from the movie.

Figure 9: Scatter plot.

and returns a DataFrame with only the instances that met this condition of 'Survived == 0'. If we want to do the inverse, we simply use != (not equals), or, alternatively == 1. We can plot these different subsets by generating to plots, storing the first in a variable, and passing that along to the second plot. Like so:

```
ax1 = died_subset.plot(x='somefeature', y='someotherfeature',
                       label='Died', color='red', kind='scatter')
ax2 = live_subset.plot(x='somefeature', y='someotherfeature',
                       label='Survived', color='green', kind='scatter',
                       ax=ax1)

ax2  # to output
```

Now notice that we're adding two new parameters here, one to label a specific subset in the legend[9], and the other to color the instances of the specific subset differently. We can also add an `alpha` parameter to determine transparency of the dots (e.g. by setting that to 0.5).

[9] We didn't have to do this before, as this was encoded in the values, here it isn't.

Task 3

- Set `Sex` to x, and `Pclass` to y. What can you conclude with this information?

- Set `Age` to x, and `Fare` to y. Is there a correlation between these two features?

- Split the above plot by `Survived`, can you see a pattern in the data points?

Testing Hypotheses

If you followed through all above tasks, you now hopefully developed some intuition regarding which features might have a relation. The Titanic was a pretty interesting reflection of society around the 1900's. It opens up some potential for statistical analyses and some richer measurements. We'll look at boxplots for this section, and some very basic pointers on how to run stats in Python.

Pandas provides the usual boxplot annotations for the mean, median, standard deviation, min, max, (if you need a refresher on these, read: link, see Figure 10) and SciPy provides a statistical test for a set variable, between the classes of the provided subgroup. Through this, you can test if for example `Fare` had an effect on your survivability. Boxplots have a `column` parameter to select features (should be in a list format), and `by` to group. So you can do something like:

```
df.boxplot(column=['Fare'], by='Survived')
```


OUTLIER More than 3/2 times of upper quartile

MAXIMUM Greatest value, excluding outliers

UPPER QUARTILE 25% of data greater than this value

MEDIAN 50% of data is greater than this value; middle of dataset

LOWER QUARTILE 25% of data less than this value

MINIMUM Least value, excluding outliers

OUTLIER Less than 3/2 times of lower quartile

Figure 10: Boxplot cheat sheet.

Then finally, `scipy.stats` also has an implementation for `ttests` and the like. Given that we have a lot of unbalanced groups and some binary data, we're not gonna go in the statistical mess required to say anything sensible about these here. Just know that the tests are available.

Task 4

- Determine a likely ticket price to survive the Titanic.

Preparing for Prediction

FROM WHAT YOU HAVE SEEN up until now, you can probably make a pretty good prediction if someone survived the Titanic. Say that we offer you this table of people as instances and some of their features:

	gender	age	pclass	fare
1	m	51	higher	80
2	f	40	middle	60
3	m	25	lower	10

How would you rank the survivability of these three persons? Now for a bit more complicated one:

	gender	age	pclass	fare
1	m	27	higher	?
2	f	21	middle	?
3	f	35	lower	?

What would you predict is (about) the fair that these three person bought their tickets for?

Data Mining

THIS IS EXACTLY THE KIND of process we want to automate using
Data Mining techniques. Up until now, we've used the Titanic data,
which is pretty simple — can be well understood in historical con-
text by humans, and therefore you can at least say something about
the likelihood of survival. The fares on the other hand require a bit
more of a complex view on the data. You can sort of guess what ball-
park the fares would be in based on the `pclass`, but the nuances are
harder. These two tasks involve prediction: survived - a discrete fea-
ture (**classification**), where we have a limited set of options to choose
from, and fare - a continuous feature (**regression**), where we'd like to
be as close to the actual number as possible.

Data Mining is in general the combination of several techniques:

1. Managing your data.

2. A thorough understanding of its contents and potential.

3. The ability to manually select and alter the data to create useful
   insights and visualizations.

4. Understanding and applying predictive models that use the full
   complexity of the data to create even better insights.

5. Making sure these models are correctly evaluated and being able
   to judge their usefulness.

6. Communicating these results.

   This course will mostly focus on 2-5. The practicals will try to
train you in 3-5. However, 2 will require your own effort but is the
most important step to the success of actually applying 3-5.

Part B - Making Predictions

For the next part, we are going to take a model-driven approach
rather than a data-driven one; we'll try to fit a Linear Regression
model (see Figure 12 for a quick visual refresher) that hopefully
predicts with low error, and see what kind of information we can get
from it.

Figure 11: Formal steps of a Data
Mining (KDD) workflow.

Figure 12: Linear Regression and
residuals (error). For 2-d (1 feature, 1
target), the fit of the line is determined
by $Y = \beta_0 + \beta_1 \cdot X$ where β_0 is the
intercept (or bias coefficient) and β_1
the coefficient for the first feature
(determining the slope of the line).

If you're still unsure how prediction and error assessment for this type of model works, try the following given bias and coefficients:

```
 -0.1084 * crime-rate +
  0.0458 * zoned +
  2.7187 * charles +
-17.376  * nitric-oxide +
  3.8016 * rooms +
 -1.4927 * employment-center +
  0.2996 * radial-highways +
 -0.0118 * property-tax +
 -0.9465 * pupil-teach-ratio +
  0.0093 * proportion-black-families +
 -0.5226 * poor-people +
 36.3411
```

You are provided with the following test data:

crime-rate	zoned	industry	charles	nitric-oxide	rooms	age	employ-ment-center	radial-highways	property-tax	pupil-teach-ratio	proportion-black-families	poor-people
0.00632	18	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98

Task 5

- Use the given coefficients and bias to predict $\hat{y}$: the median-value (by hand) for the feature vectors in the table above (X_{test}). See below for instructions.

Given the following median values (so the actual price of the houses, y):[10], use the predicted median-value to calculate the Root Mean Squared Error (RMSE) for the median-value in Footnote [10]. You do this by (for each of the feature vectors) subtracting the actual (y_i) from the predicted ($\hat{y}_i$) value, and squaring them. Here, i is the index of any instance y in Y (and thus X). After, you take the sum over all these values (should be $i = 1$ value, because 1 instance), divide it by the amount of predictions ($n = 1$), and take the root[11].

So, that was a tedious exercise you probably don't want to do again. Luckily, we can automate the fitting of the regression line, making predictions, and calculating RMSE with `scikit-learn`. In the following piece of code, we are going to recode the column names with more interpretable ones[12].

[10]

$$Y = \begin{bmatrix} 24 \end{bmatrix} \qquad (3)$$

[11]

$$\text{RMSE} = \sqrt{\frac{\sum_{i=1}^{n}(\hat{y}_i - y_i)^2}{n}} \qquad (4)$$

[12] You will also see that this is an old dataset with little regard for political correctness.

```
new_columns = ["crime-rate", "zoned", "industry", "charles",
               "nitric-oxide", "rooms", "age", "employment-center",
               "radial-highways", "property-tax", "pupil-teach-ratio",
               "proportion-black-families", "poor-people", "median-value"]

df2 = pd.DataFrame.from_csv('./data/Housing/housing.csv', index_col=None)

df2.columns = new_columns  # replace the initial columns with new names
df2
```

Now, we want to split `median-value` into its own variable. We can `pop` it off the dataframe, and store it in a new variable (as a list):

```
y = df2.pop("median-value").tolist()
```

Now because we removed `median-value` from `df2`, we can immediately convert it to a matrix to get X. As such:

```
X = df2.as_matrix()
```

Now, we want to create a train and test set. We can do this pretty easily with the following piece of code:

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y)
```

Here, `train_test_split` returns multiple objects, so that's why there's more variable names before the = sign. For regression, we first initialize the classifier (this is when you'd want to submit extra) parameters. After, we can call the `fit()` method on our training data, and labels[13].

```
from sklearn.linear_model import LinearRegression

lr = LinearRegression()
lr.fit(X_train, y_train)
```

The classifier is now fitted, or trained. We can first try to interpret the coefficients. This can be done by calling the `coef_` attribute, and then zipping (concatenating)[14] it with the column names we defined before (otherwise, you'd just call `df2.columns`), like below:

```
list(zip(columns[:-1], lr.coef_))
```

Now, we can also do prediction on our **test** set, and then evaluate our prediction compared to the *actual* true labels:

[13] `scikit-learn` has very comprehensible documentation. Look up the parameters for the model and its methods!

[14] No need to really understand what's going on here, just remember it as a way to look at the coefficients.

```
from sklearn.metrics import mean_squared_error
import numpy as np

y_pred = lr.predict(X_test)
np.sqrt(mean_squared_error(y_test, mean_squared_error))
```

Then finally, solve:

Task 5

- Fit Linear Regression on the Boston Housing dataset, interpret the coefficients. What would these mean?

- How do we asses if the model has a) actually learned something, and b) will generalize well?

Solutions

Task 1:

- Given the Titanic Wiki quote, we see that there were 1317 passengers. So we're missing $1317 - 891$ passengers (see `df.count()`).

- This is personal: if you aren't sure, look at the docs what they mean, and **especially** what the values are!

- Well, if you read the quote carefully we know that it's highly unlikely that there's any crew in the dataset. Some features could however be: Fare (if that would be 0), or Embarked (see Wiki).

Task 2:

- `Pclass` and `Sex`.

- Women and children first, ice hit the boat in the lower deck (lower class). Most of the latter drowned before even making it to any lifeboats.

Task 3:

- Not more than that they are two discrete features and scatter plots only work for continuous ones[15].

  [15] See first bit of code below to fix the error pandas throws.

- Not strong enough to clearly see a visual correlation. You can test it using `scipy.stats.pearsonr`. See below for code.

- There's a few observations we can make, but not necessarily very clear patterns. It looks like from low-fare payers over 40, the majority seems to be a non-survivor. We can also see that the high fare payers generally survived (there's even a few super high outliers with 500+ fare, they survived).

```
df['Sex'] = pd.get_dummies(df['Sex'])  # convert numerical to discrete

from scipy.stats import pearsonr

pearsonr(df['Age'], df['Fare'])  # run df['Age'].fillna(0) before!
```

Task 4:

- With paid Fair upwards from 50, you're an outlier if you die.

Task 5:

- Positive means it increases the house price by some factor, negative coefficients mean it decreases the price by this factor. Note that they are on different scales though!

- a) we would need some 'stupid' baseline that we can use to predict without learning anything (like the mean of all the values that we saw in our training dataset). b) we have a test set, so we know the performance on unseen data. Actual generalizability isn't easy to comment on with small datasets. The more instances we have (especially across time, these are housing market dynamics mind you), the more confident we can be that it still holds.